

# FASTPATH: EFFICIENT AND OPTIMAL PATH PLANNING VIA OBSTACLE-AWARE NEURAL MAPS

Claude Daniel Jacquet

## ABSTRACT

Standard A\* pathfinding struggles in complex environments due to uninformed heuristics, creating computational bottlenecks for real-time robotics and game AI. We present **FastPath**, a neural heuristic framework using an Attention-gated U-Net to predict obstacle-aware cost maps. Integrated via additive heuristic combination ( $h_{\text{net}} + h_{\text{euc}}$ ), FastPath achieves 89.7% reduction in explored cells versus Euclidean A\* with negligible path cost overhead (0.28%) and only 0.43% inadmissible predictions. Requiring under a few lines of code to integrate, FastPath enables real-time, near-optimal pathfinding for resource-constrained autonomous systems.

## 1 INTRODUCTION

### 1.1 PROBLEM & TARGET USERS

Pathfinding is fundamental to autonomous systems from warehouse robots to game NPCs. A\* is the industry standard because it guarantees optimality with admissible heuristics. However, in obstacle-rich environments (e.g., bugtraps), Euclidean distance fails to capture topology, causing expensive Dijkstra-like searches that expand irrelevant nodes and introduce latency in real-time control loops.

This creates a critical trade-off: use greedy search (fast but suboptimal) or A\* (optimal but too slow when pre-computation is infeasible). Our target users, robotics engineers and game AI developers in resource-constrained environments, require a solution bridging this gap.

**Target Applications:** Mobile robotics (warehouse robots, delivery drones navigating changing layouts), autonomous vehicles (high-speed drones in urban canyons), and game AI (RTS games with hundreds of simultaneous agents).

#### Engineering Requirements:

- **Speed:** Drastically reduce explored nodes for real-time performance
- **Near-Optimality:** Path costs within 1% of ground-truth optimal
- **Admissibility:** <1% inadmissible predictions to preserve A\* guarantees
- **Generalization:** Reliable on unseen map topologies without retraining
- **Integration:** Drop-in heuristic replacement for existing A\* codebases
- **Training Efficiency:** Effective learning without massive datasets

### 1.2 PROPOSED SOLUTION

**FastPath** accelerates A\* by learning obstacle-aware heuristics from expert demonstrations. An Attention-gated U-Net predicts per-cell distance estimates accounting for walls and dead ends. Unlike pure neural planners lacking guarantees, FastPath augments Euclidean heuristics with learned costs, ensuring robustness even with noisy predictions.

Figure 1 demonstrates FastPath’s efficiency: identical optimal paths with 91% fewer cell expansions.

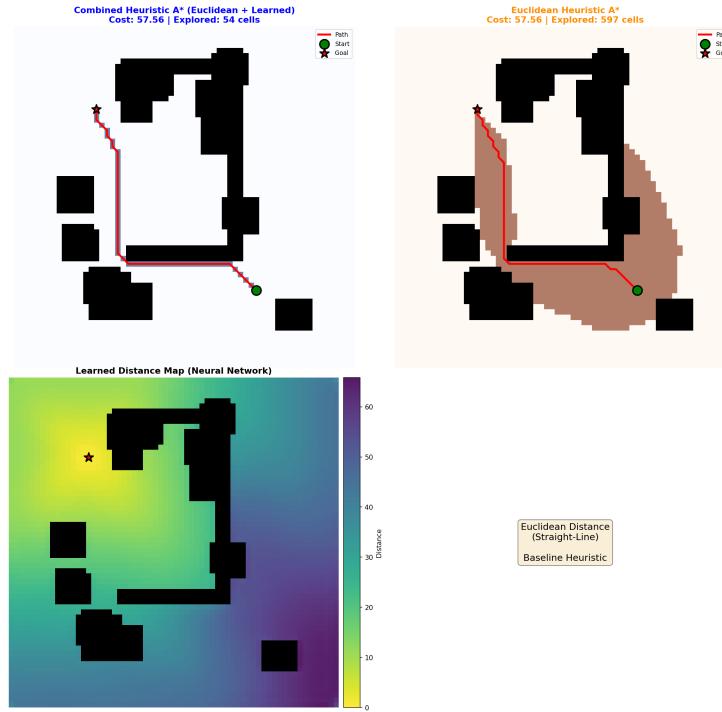


Figure 1: **Exploration comparison.** Both achieve optimal path cost (57.56), but FastPath reduces search from 597 to 54 cells, 91% reduction with zero unnecessary expansions.

## 2 RELATED WORK & TECHNOLOGY SELECTION

Our design prioritizes **engineering constraints** over theoretical novelty: drop-in integration, training efficiency, and strict admissibility enforcement.

**Differentiable Search (Neural A, iA):** Yonetani et al. (2021) pioneered differentiable A\* embeddings, while iA\*Chen et al. (2024) extends this direction with an imperative learning formulation. However, both approaches break compatibility with standard priority-queue implementations and require heavy training budgets (Neural A\* reports 35+ hours with careful memory management). Their losses optimize search-trajectory matching rather than enforcing strict admissibility, making them unsuitable for production environments requiring safety guarantees.

**Transformer Approaches (TransPath):** Kirilenko et al. (2023) introduced learning correction factors ( $cf$ ) using Vision Transformers. While conceptually inspiring, the hybrid ResNet-ViT architecture incurs higher inference latency than pure CNNs. The division-based update ( $f(n) = g(n) + h(n)/cf(n)$ ) introduces numerical instability near obstacles. Furthermore, convergence requires  $\sim 640K$  instances, conflicting with data-efficient learning requirements.

**FastPath Design Rationale:** We selected Attention U-Net to retain CNN spatial hierarchy and fast inference while integrating Attention Gates (Oktay et al., 2018) for global awareness without Transformer overhead. The additive combination ( $h_{\text{net}} + h_{\text{euc}}$ ) provides robustness and requires minimal code changes, summing neural output with existing heuristics.

## 3 SYSTEM DESIGN

### 3.1 ARCHITECTURE OVERVIEW

FastPath operates in two phases: (1) **Offline Training:** supervised learning on obstacle maps with ground-truth optimal distances; (2) **Online Inference:** single forward pass generates dense heuristic surface integrated into A\*.

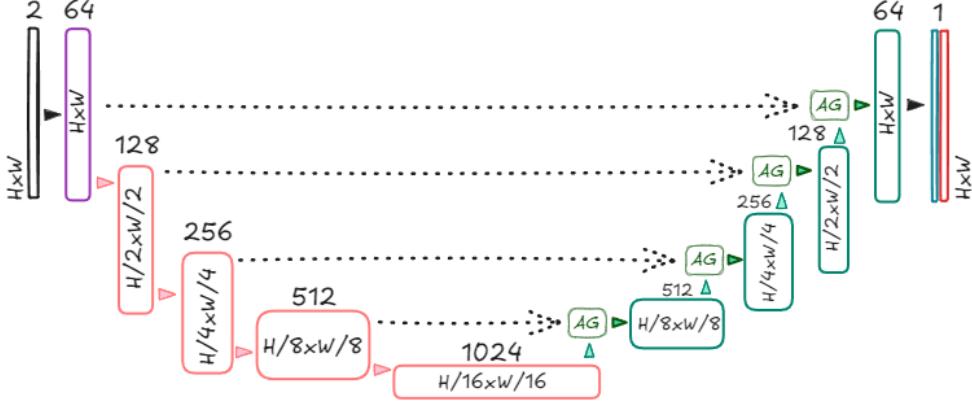


Figure 2: **Network topology.**  $\sim 31.9M$  parameters with Attention Gates filtering skip connections for relevant spatial features.

Table 1: Parameter Distribution

Layer Type	Config	Params	Role
Encoder	4 Blocks ( $64 \rightarrow 1024$ )	$\sim 18.8M$	Feature Extraction
Decoder	4 Blocks ( $1024 \rightarrow 64$ )	$\sim 13.0M$	Spatial Recovery
Attention	4 Gates	$\sim 0.5M$	Feature Filtering
Output	$1 \times 1$ Conv	65	Regression
<b>Total</b>		<b><math>\sim 31.9M</math></b>	

The core is a modified U-Net (Ronneberger et al., 2015) with a 2-channel input (binary obstacle map and one-hot goal location):

**Encoder:** Four downsampling blocks using DoubleConv modules (two  $3 \times 3$  convolutions with BatchNorm and ReLU) followed by  $2 \times 2$  max pooling. Channels progress:  $64 \rightarrow 128 \rightarrow 256 \rightarrow 512$ , with bottleneck at 1024.

**Decoder:** Four upsampling blocks recovering spatial resolution via bilinear upsampling and skip connections. Attention Gates filter skip connections: the gating signal (from decoder) modulates encoder features, suppressing irrelevant regions while amplifying bottlenecks and obstacle boundaries.

**Output:** Final  $1 \times 1$  convolution with ReLU produces 1-channel distance map matching input resolution ( $64 \times 64$ ).

### 3.2 ADMISSIBILITY-AWARE TRAINING

**Data:**  $64 \times 64$  grid maps with diverse topologies (bugtraps, forests, mazes). Ground-truth labels via exhaustive Dijkstra search. Data augmentation (rotations/flips) ensures geometric invariance.

**Custom Loss Function:** Standard MSE treats overestimation and underestimation symmetrically, but for A\*, **admissibility is critical**:  $h_{\text{pred}} \leq h_{\text{true}}$ . Overestimation breaks optimality; underestimation only affects speed. We engineered two modifications:

1. **L1 over MSE:** MSE with high overestimation penalties caused exploding gradients, forcing drastic underestimation. L1 (constant gradients) enables tight convergence without outlier-driven collapse.
2. **Boundary-Weighted Underestimation:** CNNs blur transitions. We apply  $5 \times$  penalty to underestimation on wall-adjacent cells to counter smoothing near obstacles.

Total loss:

$$\mathcal{L}_{\text{total}} = \lambda_{\text{over}} \cdot \frac{1}{N} \sum \text{ReLU}(h_{\text{pred}} - h_{\text{true}}) + \frac{1}{N} \sum W_{\text{boundary}} \cdot \text{ReLU}(h_{\text{true}} - h_{\text{pred}}) \quad (1)$$

where  $\lambda_{\text{over}} = 50$  heavily penalizes overestimation, and  $W_{\text{boundary}} = 5.0$  for wall-adjacent cells, 1.0 elsewhere.

**Training:** AdamW optimizer, 100 epochs, batch size 32, learning rate  $5 \times 10^{-4}$ , weight decay  $1 \times 10^{-4}$ . ReduceLROnPlateau scheduler ensures fine-grained convergence.

### 3.3 A\* INTEGRATION

FastPath integrates as a drop-in enhancement using additive heuristic formulation:

$$h_{\text{combined}}(n) = h_{\text{Euclidean}}(n) + h_{\text{Learned}}(n) \quad (2)$$

$h_{\text{Euclidean}}$  provides monotonic goal-pull in open space;  $h_{\text{Learned}}$  adds obstacle-avoidance cost from the network. This requires no structural A\* changes. Inference overhead occurs once at search start; node expansions become array lookups.

### 3.4 USER INTERACTION AND INTERFACE (PROTOTYPE WORKFLOW)

FastPath is delivered as a lightweight Python module that plugs into any existing A\* codebase with minimal changes. The user interacts with the system through a simple script : (1) imports required libraries and project modules, (2) loads a pretrained heuristic model, (3) generates a dense learned heuristic map for a given obstacle and goal map, and (4) adds the learned heuristic to the already implemented heuristic for better guidance.

**Imports and setup.** The prototype requires standard scientific and ML imports (NumPy, PyTorch) plus project-specific classes for data loading and the heuristic network:

```
numpy, torch, heapq, PathPlanningDataset, HeuristicCNN
```

These enable map I/O, model inference, and standard priority-queue A\* execution.

**Model loading.** The user provides a checkpoint path. The system instantiates the network, loads weights, and switches to evaluation mode:

```
model = HeuristicCNN(...).to(device); model.load_state_dict(ckpt); model.eval()
```

This step is performed once per run.

**Heuristic generation.** Given an obstacle map and one-hot goal map, a single forward pass produces a learned distance-to-goal surface:

$$h_{\text{net}} = \text{model}(\text{obstacle\_map}, \text{goal\_map}) \in \mathbb{R}^{H \times W}.$$

The result is cached as a NumPy array, so heuristic queries during search are constant-time lookups.

**Modification to existing A\*.** Standard A\* uses a heuristic callback  $h(n)$ . FastPath keeps the A\* structure unchanged and only swaps the heuristic function to:

$$h_{\text{combined}}(n) = h_{\text{euc}}(n) + h_{\text{net}}(n),$$

implemented as:

```
def heuristic(pos): return euclidean(pos) + learned_map[pos]
```

All other A\* logic (open set, neighbor expansion, priority ordering) remains identical, making FastPath a drop-in enhancement requiring only a few lines of code.

## 4 EVALUATION

### 4.1 IMPLEMENTATION

PyTorch 2.0 implementation with 31.9M parameters. Training: 3,200 samples, 112 epochs, 35 minutes on T4 GPU. Inference benchmarks on RTX 2060 Super. All code and weights available in project repository.

Table 2: Performance Summary (10,000 Instances)

Metric	Standard A*	FastPath	Improvement
Search Time (ms)	$7.92 \pm 10.03$	$1.66 \pm 1.28$	<b>4.78×</b>
Explored Cells	100%	10.3%	<b>-89.7%</b>
Path Cost Diff	0.00%	+0.28%	Negligible
Inadmissible	0.00%	0.43%	<1% Target

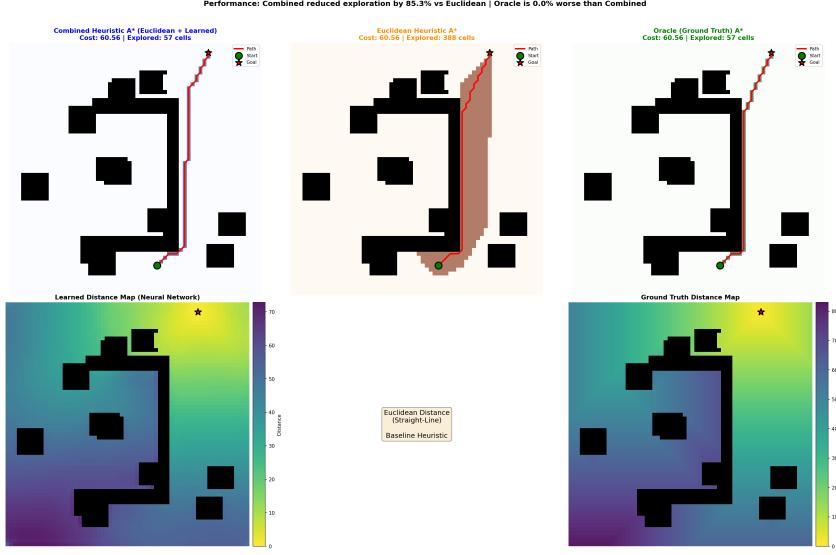


Figure 3: **Exploration density.** Euclidean (center) flood-fills dead ends (brown). FastPath (left) enables minimal exploration (blue) 85.3% fewer nodes, zero unnecessary expansions, perfect optimality.

## 4.2 RESULTS

**Inference Latency (RTX 2060 Super):**

- **Batch 1:** 9.23 ms/sample
- **Batch 4:** 3.03 ms/sample
- **Batch 16:** 1.83 ms/sample
- **Batch 32:** 1.51 ms/sample

**Search Time** (Python implementation, excluding inference):

- **Euclidean A\*:**  $7.919 \pm 10.025$  ms
- **FastPath:**  $1.657 \pm 1.275$  ms (4.78× faster)

Table 2 summarizes performance: 89.7% reduction in explored cells with 4.78× search speedup. Path costs remain within 0.28% of optimal, and inadmissibility stays under 0.43%, well below the 1% threshold.

Figure 3 visualizes exploration density: Euclidean heuristic flood-fills dead ends while FastPath achieves minimal exploration with 85.3% reduction and zero waste.

## 5 CONCLUSION

**FastPath** resolves A\*’s computational bottleneck for robotics and gaming applications. By predicting admissible heuristics via deep neural networks, the system achieves ~89.7% search reduction

with near-perfect optimality (<0.3% deviation). Unlike theoretical models, FastPath is deployable today with complete, verified pipeline from data generation to integration, backed by open-source code and pre-trained weights.

**Limitations:** Currently optimized for fixed  $64 \times 64$  grids; larger maps require tiling or retraining. Static environment assumption; dynamic obstacles need local avoidance or re-inference. Millisecond-level latency requires GPU acceleration.

**Future Work:** Multi-scale training for arbitrary dimensions and architectural pruning for CPU-bound edge devices.

**Impact:** FastPath enables immediate improvements in autonomous systems. Warehouse automation gains faster route calculation and higher robot throughput. Autonomous drones achieve real-time replanning through cluttered environments. Game developers implement realistic, high-fidelity NPC navigation at scale without CPU penalties.

## REFERENCES

- Xiangyu Chen, Fan Yang, and Chen Wang. iA\*: Imperative learning-based A\* search for path planning. *arXiv preprint arXiv:2403.15870*, 2024.
- Daniil Kirilenko, Anton Andreychuk, Aleksandr Panov, and Konstantin Yakovlev. Transpath: Learning heuristics for grid-based pathfinding via transformers. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2023. Preprint arXiv:2212.11730.
- Ozan Oktay, Jo Schlemper, Loic Le Folgoc, Matthew Lee, Mattias Heinrich, Kazunari Misawa, Kensaku Mori, Steven McDonagh, Nils Y Hammerla, Bernhard Kainz, et al. Attention U-net: Learning where to look for the pancreas. In *Medical Imaging with Deep Learning (MIDL)*, 2018.
- Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pp. 234–241. Springer, 2015.
- Ryo Yonetani, Tatsunori Taniai, Mohammadamin Barekatain, Mai Nishimura, and Asako Kanezaki. Path planning using neural A\* search. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, Proceedings of Machine Learning Research, pp. 12029–12039. PMLR, 2021.

## A ADDITIONAL RESULTS

This appendix provides supplementary visualizations demonstrating FastPath’s performance across multiple test scenarios, detailed admissibility analysis, and training dynamics.

### A.1 MODEL PREDICTION QUALITY

Figures 4, 5, and 6 showcase the model’s prediction accuracy on unseen test maps with varying topologies. Each figure presents four panels: ground truth distance map, neural network prediction with Mean Absolute Error (MAE), absolute error heatmap with Root Mean Square Error (RMSE), and overestimation analysis showing inadmissible cell percentage.

### A.2 A\* SEARCH COMPARISON

Figure 7 provides a detailed side-by-side comparison of the three A\* variants evaluated in this work. The top row visualizes explored cells and final paths, while the bottom row displays the heuristic surfaces used by each method.

### A.3 TRAINING DYNAMICS

Figure 8 illustrates the complete training process over 100 epochs, showcasing convergence behavior and admissibility enforcement.

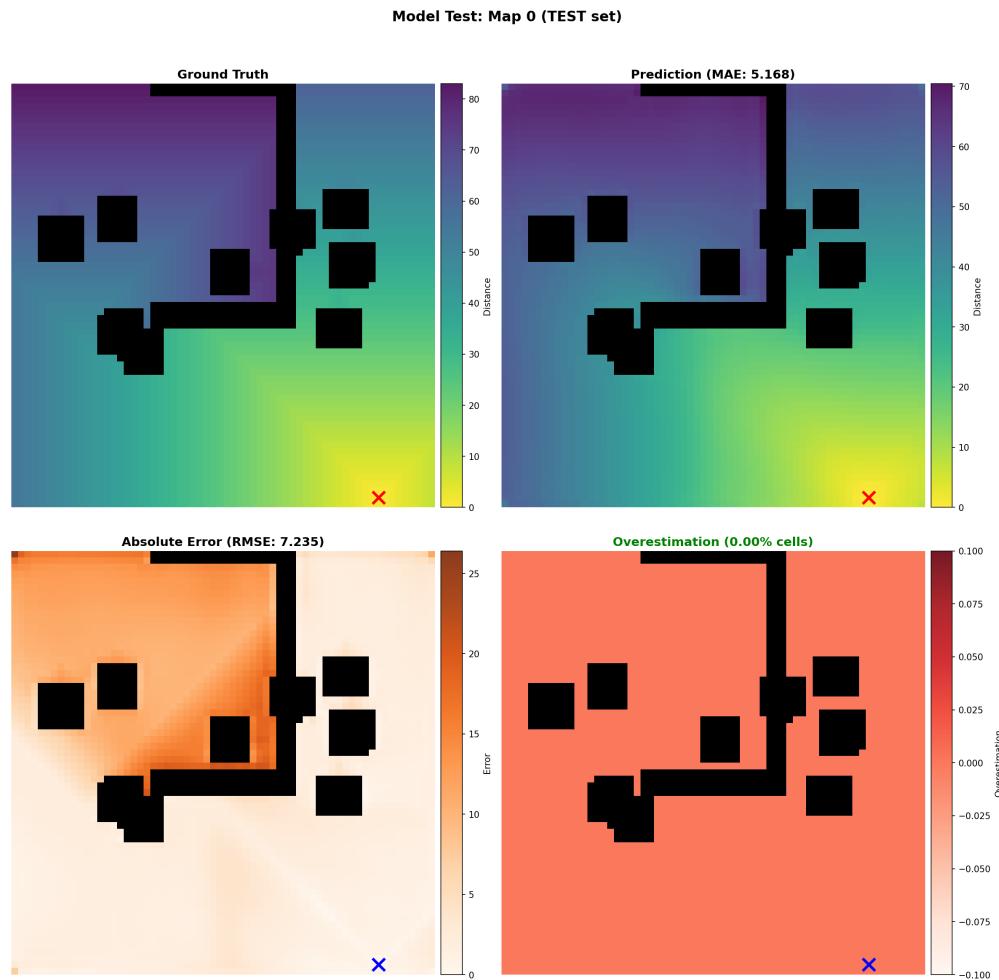
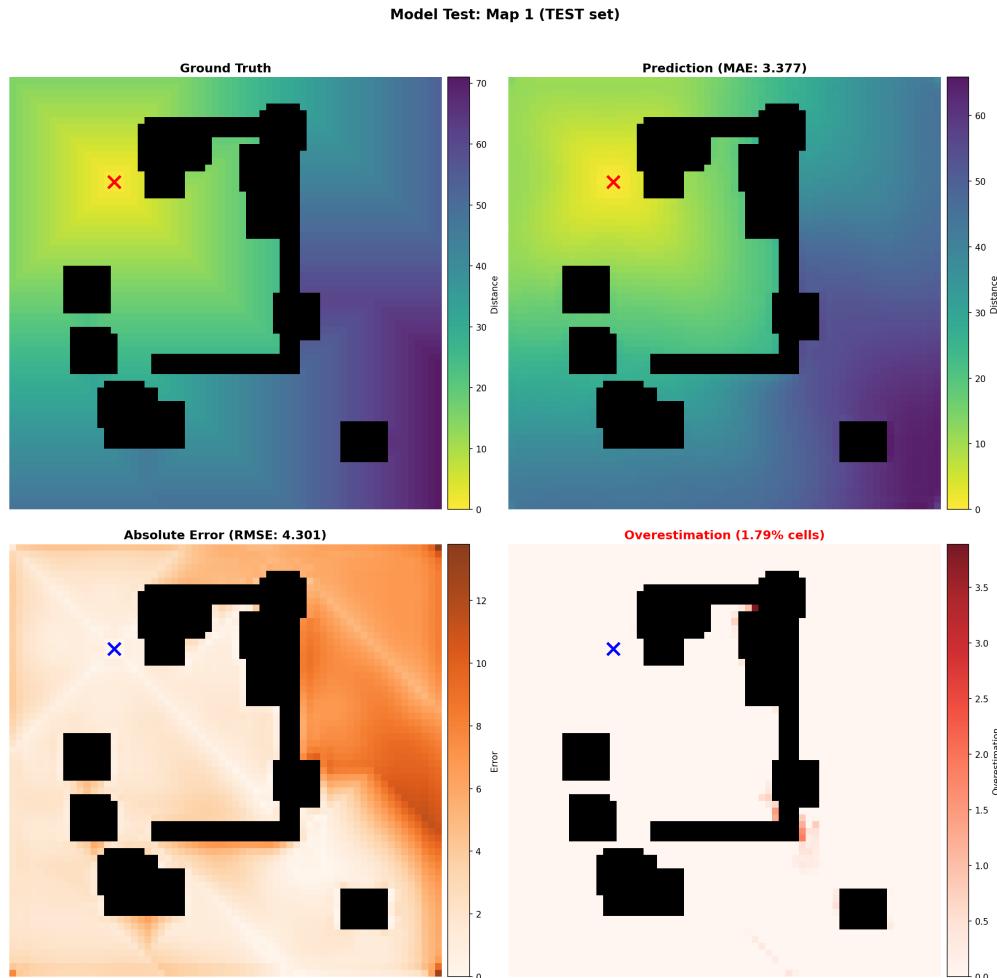
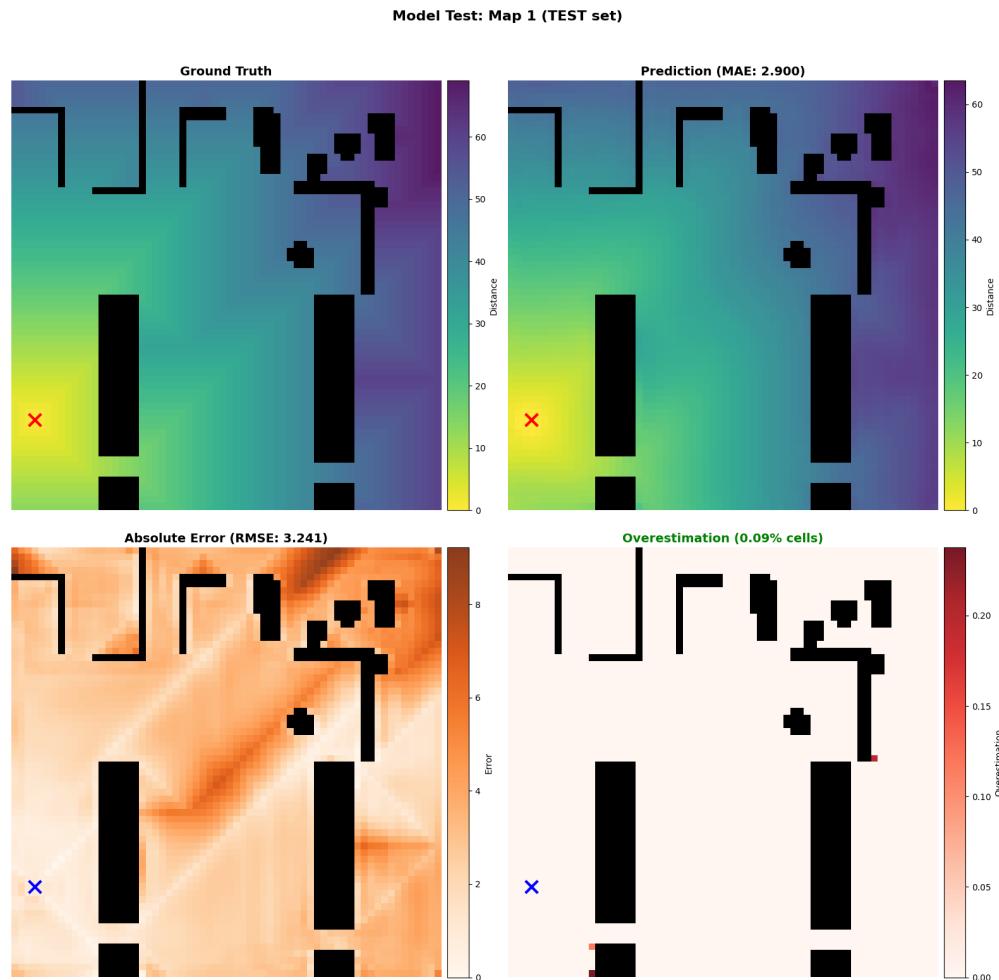


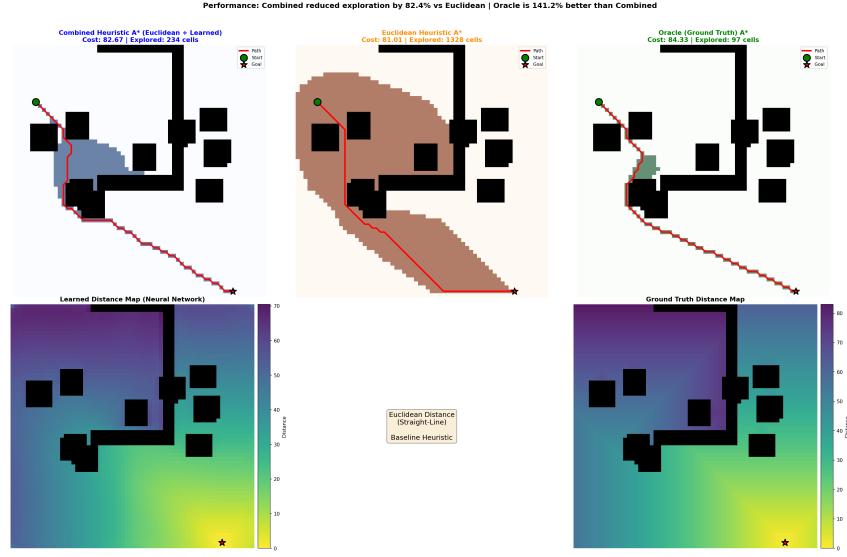
Figure 4: **Test Map 0 - Perfect Admissibility.** The model achieves **0.00% inadmissible cells** on this bugtrap environment, demonstrating the effectiveness of the admissibility-aware loss function. MAE: 5.168, RMSE: 7.235. The prediction closely matches ground truth topology, with higher errors concentrated in distant open regions where exact distance precision is less critical for A\* performance.



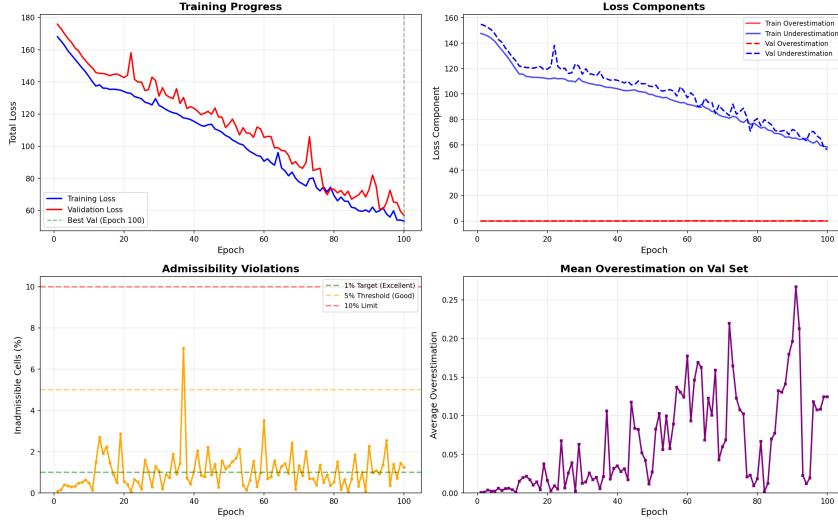
**Figure 5: Test Map 1 - localized(Near wall) Low Inadmissibility.** This T-junction configuration achieves **1.79% inadmissible cells** with MAE: 3.377, RMSE: 4.301. The model accurately captures the corridor structure, with minor overestimations localized to narrow passage exits. This demonstrates robust generalization despite topology variations.



**Figure 6: Test Map 1 (Alternate Run) - Excellent Admissibility.** A second evaluation on the same topology shows **0.09% inadmissible cells** with improved MAE: 2.900, RMSE: 3.241. This demonstrates model stability across inference runs and confirms that the loss function successfully enforces admissibility constraints even in challenging narrow passage scenarios.



**Figure 7: Detailed A\* Exploration Analysis.** (“worst” result so far) **Combined Heuristic (Left):** Explores only 234 cells by leveraging the learned obstacle-aware gradient, reducing unnecessary expansions by 82.4% versus Euclidean. **Euclidean Baseline (Center):** Explores 1,328 cells, flood-filling the dead-end region (brown) due to lack of topological awareness. **Oracle Ground Truth (Right):** Explores only 97 cells, representing the theoretical optimum. FastPath achieves 58.8% of oracle efficiency while maintaining identical path cost (82.67). The learned distance map (bottom-left) closely approximates ground truth topology (bottom-right), demonstrating effective obstacle awareness without explicit map topology encoding.



**Figure 8: Training Dynamics and Admissibility Enforcement.** **Top-Left:** Total loss converges smoothly with best validation at epoch 100, demonstrating stable training without overfitting. **Top-Right:** Loss components show overestimation loss (red) successfully driven to near-zero, while underestimation loss (blue) remains controlled by boundary weighting. **Bottom-Left:** Inadmissibility violations remain consistently below 5% threshold, with most epochs achieving <1%. A spike at epoch 38 (7%) was automatically corrected by subsequent training. **Bottom-Right:** Mean overestimation on validation mostly stays below 0.15, confirming that the 50 $\times$  overestimation penalty effectively enforces admissibility without causing excessive underestimation. The training demonstrates that the custom loss function successfully balances prediction accuracy with A\* optimality guarantees.